

Quora_Data_Challenge

November 3, 2019

1 Quora Data Challenge

Suppose you are a Data Scientist on the Mobile team at Quora. The team has just introduced a new UI design to the Quora app. The goal of the new design is to increase user engagement (measured by minutes spent on site). The team ran an A/B test to evaluate the change. Using the data, help the team understand the impact of the UI change better.

```
In [1]: import pandas as pd
import datetime
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [7]: #take a look at different tables
#This table contains active minutes data logged after experiment started.
#Each row represents the total number of minutes spent on site for each user on a date.
#If a user never visited the site for a given date, there wouldn't be data for that uid
t1 = pd.read_csv("t1_user_active_min.csv")
print("the number of records is ",t1.shape)
t1.head()
t1['active_mins'].describe()
```

```
the number of records is (1066402, 3)
```

```
Out[7]: count    1.066402e+06
mean      3.616809e+01
std       1.270484e+03
min       1.000000e+00
25%       2.000000e+00
50%       5.000000e+00
75%       1.700000e+01
max       9.999900e+04
Name: active_mins, dtype: float64
```

```
In [49]: len(t1.uid.unique())
```


```
Out[49]: 46525
```

```
In [10]: # By definition, a user cannot be more than 24*60 minutes using quora on a given day.
# I assume these are some hard_coded value that has some meaning,e.g., for some reason,
# Quora could not capture the number of minutes on the given day.
t1[(t1.active_mins > (24*60))]
# We have 172 records where active_mins >1440 which contradicts the metrics definition.
# I will filter them out.
t1 = t1[(t1.active_mins<=1440)]
t1.shape
```

```
Out[10]: (1066230, 3)
```


```
In [12]: %matplotlib inline
t1['active_mins'].plot.hist()
t1['active_mins'].describe()
```

```
Out[12]: count      1.066230e+06
mean        2.004248e+01
std         4.653763e+01
min         1.000000e+00
25%         2.000000e+00
50%         5.000000e+00
75%        1.700000e+01
max         8.970000e+02
Name: active_mins, dtype: float64
```




```
In [13]: sns.boxplot(x=t1['active_mins'])
# we still observe some outliers. I will explore the possibility of filtering the out.
```

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66beb1b00>



```
In [15]: t1=t1[(t1.active_mins<(20+(3*46)))]  
sns.boxplot(x=t1['active_mins'])
```

Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66be7f390>




```
In [22]: # compute IQ range to filter outliers
q1 = t1['active_mins'].quantile(0.25)
q3 = t1['active_mins'].quantile(0.75)
print(q1,q3)
iqr = q3 - q1
cutoff = 1.5*iqr
t1 = t1[((t1.active_mins)> q1 - cutoff)]
t1 = t1[((t1.active_mins)< q3 + cutoff)]
```

2.0 9.0

Out[22]: (828130, 3)

```
In [23]: t1.shape
sns.boxplot(x=t1['active_mins'])
```

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66be015c0>



```
In [4]: #This table contains users treatment assignment.
        #Each row represents the assignment information for a unique user.
        t2 = pd.read_csv("t2_user_variant.csv")
        print("the number of records is ",t2.shape)
        t2.head()
```

the number of records is (50000, 4)

```
Out[4]:
```

	uid	variant_number	dt	signup_date
0	0	0	2019-02-06	2018-09-24
1	1	0	2019-02-06	2016-11-07
2	2	0	2019-02-06	2018-09-17
3	3	0	2019-02-06	2018-03-04
4	4	0	2019-02-06	2017-03-09

```
In [26]: # check that we have a unique value for dt.
        t2.dt.unique()
```

```
Out[26]: array(['2019-02-06'], dtype=object)
```

```
In [162]: t2.variant_number.value_counts()
```

```
Out[162]: 0    40000
          1    10000
          Name: variant_number, dtype: int64
```

```

In [24]: #join by userid
merged_df = pd.merge(left=t1,right=t2, left_on='uid', right_on='uid')

In [48]: len(merged_df.uid.unique())

Out[48]: 46525

In [ ]: merged_df.shape

In [42]: merged_df.head(15)

Out[42]:
```

	uid	dt_x	active_mins	variant_number	dt_y	signup_date
0	0	2019-02-22	5.0	0	2019-02-06	2018-09-24
1	0	2019-03-11	5.0	0	2019-02-06	2018-09-24
2	0	2019-03-18	3.0	0	2019-02-06	2018-09-24
3	0	2019-03-22	4.0	0	2019-02-06	2018-09-24
4	0	2019-04-03	9.0	0	2019-02-06	2018-09-24
5	0	2019-04-06	1.0	0	2019-02-06	2018-09-24
6	0	2019-04-17	1.0	0	2019-02-06	2018-09-24
7	0	2019-05-07	3.0	0	2019-02-06	2018-09-24
8	0	2019-05-14	1.0	0	2019-02-06	2018-09-24
9	0	2019-05-19	1.0	0	2019-02-06	2018-09-24
10	0	2019-05-22	3.0	0	2019-02-06	2018-09-24
11	0	2019-06-14	5.0	0	2019-02-06	2018-09-24
12	0	2019-06-16	2.0	0	2019-02-06	2018-09-24
13	1	2019-02-16	13.0	0	2019-02-06	2016-11-07
14	1	2019-03-17	5.0	0	2019-02-06	2016-11-07

```

In [32]: # double check that a user is only in one group
aux=merged_df.groupby('uid').variant_number.nunique()
max(aux)

Out[32]: 1

In [38]: merged_df.dtypes
min_date = min(merged_df['dt_x'])
max_date = max(merged_df['dt_x'])
print('min date is',min_date)
print('max date is',max_date)

min date is 2019-02-06
max date is 2019-07-05

In [40]: # count records in each group
merged_df.variant_number.value_counts()

Out[40]: 0    693417
         1    134713
         Name: variant_number, dtype: int64

```

```
In [46]: # group by user to compute mean time across all his/her visit.
df = merged_df.groupby(['uid', 'variant_number'], as_index=False)['active_mins'].mean()
df.shape # we have 46525 unique users in both groups
df.variant_number.value_counts()
```

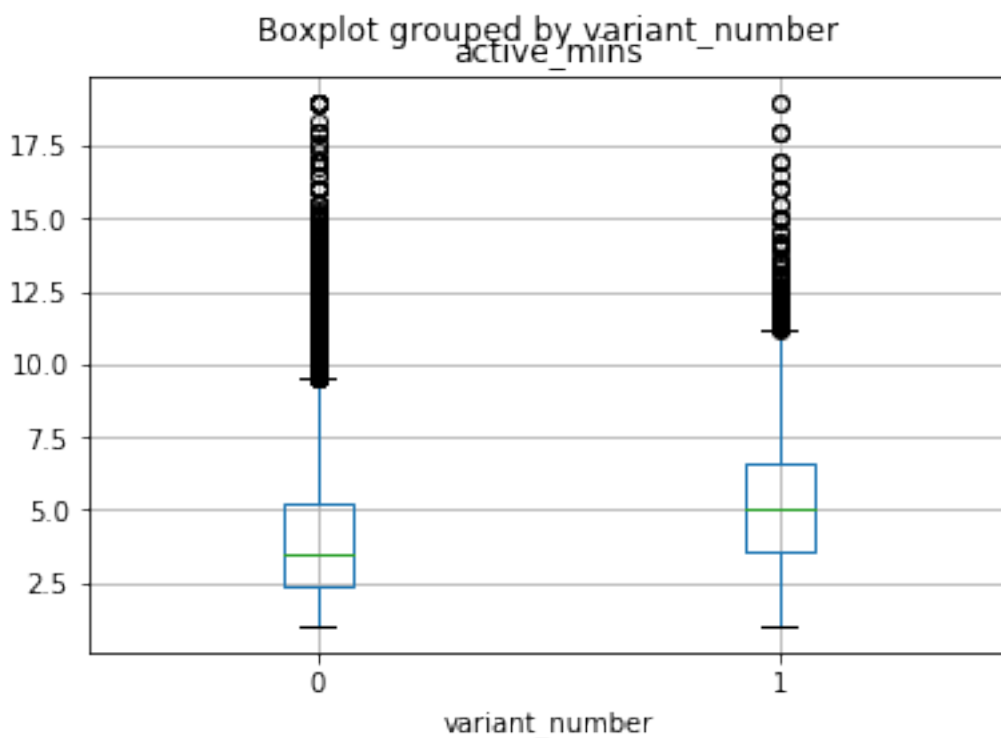
```
Out[46]: 0    37362
         1     9163
         Name: variant_number, dtype: int64
```

```
In [53]: df.head(10)
df.groupby('variant_number', as_index=False)['active_mins'].mean()
```

```
Out[53]:   variant_number  active_mins
         0                0    4.149859
         1                1    5.291342
```

```
In [51]: %matplotlib inline
df.boxplot(column='active_mins', by='variant_number')
```

```
Out[51]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66bcbd3c8>
```



```
In [58]: # compute confidence interval on difference of means.
stats = df.groupby(['variant_number'])['active_mins'].agg(['mean', 'count', 'std', 'var'])
stats
```

```
Out[58]:
```

	mean	count	std	var
variant_number				
0	4.149859	37362	2.40873	5.801981
1	5.291342	9163	2.34741	5.510334

```
In [60]: stats.loc[0]['var']
```

```
Out[60]: 5.80198115819105
```

```
In [65]: import math
# estimated standard deviation of the difference between the two means.
sigma_dif = math.sqrt((stats.loc[0]['var']/stats.loc[0]['count'])+((stats.loc[1]['var']
upper = (stats.loc[1]['mean']-stats.loc[0]['mean']) + (1.96 * sigma_dif)
lower = (stats.loc[1]['mean']-stats.loc[0]['mean']) - (1.96 * sigma_dif)
print([lower,upper])
```

```
[1.08756830375531, 1.1953974370435352]
```

```
In [67]: #This table contains active minutes data before the experiment started.
#It has a similar format as t1, except the dt range can extend before
#the experiment start date.
t3 = pd.read_csv("t3_user_active_min_pre.csv")
print("the number of records is ",t3.shape)
t3.head()
```

```
the number of records is (1190093, 3)
```


```
Out[67]:
```

	uid	dt	active_mins
0	0	2018-09-24	3.0
1	0	2018-11-08	4.0
2	0	2018-11-24	3.0
3	0	2018-11-28	6.0
4	0	2018-12-02	6.0

```
In [70]: t3 = t3[(t3.active_mins<=1440)]
%matplotlib inline
t3['active_mins'].plot.hist()
t3['active_mins'].describe()
```


```
Out[70]:
```

count	1.189927e+06
mean	1.825735e+01
std	4.354546e+01
min	1.000000e+00
25%	2.000000e+00
50%	4.000000e+00
75%	1.400000e+01
max	8.970000e+02
Name: active_mins, dtype: float64	




```
In [71]: sns.boxplot(x=t3['active_mins'])
```

```
Out[71]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66bc438d0>
```



```
In [74]: t3=t3[(t3.active_mins<(18.25+(3*43.54)))]
sns.boxplot(x=t3['active_mins'])
```

```
Out[74]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66bb46d30>
```




```
In [75]: # compute IQ range to filter outliers
q1 = t3['active_mins'].quantile(0.25)
q3 = t3['active_mins'].quantile(0.75)
print(q1,q3)
iqr = q3 - q1
cutoff = 1.5*iqr
t3 = t3[((t3.active_mins)> q1 - cutoff)]
t3 = t3[((t3.active_mins)< q3 + cutoff)]
```

```
2.0 8.0
```

```
In [76]: sns.boxplot(x=t3['active_mins'])
```

```
Out[76]: <matplotlib.axes._subplots.AxesSubplot at 0x7fc66b759390>
```



```
In [80]: #check the dates of dataframe
min_date = min(t3['dt'])
max_date = max(t3['dt'])
print('min_date is',min_date)
print('max_date is',max_date)
```

```
min_date is 2018-08-10
max_date is 2019-02-05
```

```
In [81]: merged_df2 = pd.merge(left=t3,right=t2, left_on='uid', right_on='uid')
```

```
In [82]: merged_df2.head(10)
```

```
Out[82]:
```

	uid	dt_x	active_mins	variant_number	dt_y	signup_date
0	0	2018-09-24	3.0	0	2019-02-06	2018-09-24
1	0	2018-11-08	4.0	0	2019-02-06	2018-09-24
2	0	2018-11-24	3.0	0	2019-02-06	2018-09-24
3	0	2018-11-28	6.0	0	2019-02-06	2018-09-24
4	0	2018-12-02	6.0	0	2019-02-06	2018-09-24
5	0	2018-12-04	1.0	0	2019-02-06	2018-09-24
6	0	2018-12-07	8.0	0	2019-02-06	2018-09-24
7	0	2018-12-09	5.0	0	2019-02-06	2018-09-24
8	0	2018-12-14	8.0	0	2019-02-06	2018-09-24
9	0	2018-12-15	2.0	0	2019-02-06	2018-09-24

```

In [84]: merged_df2.variant_number.unique()
merged_df2.variant_number.value_counts()

Out[84]: 0    754497
         1    164195
         Name: variant_number, dtype: int64

In [85]: df2 = merged_df2.groupby(['uid', 'variant_number'], as_index=False)['active_mins'].mean()
df2.head()

Out[85]:   uid  variant_number  active_mins
0     0                0      3.333333
1     1                0     11.333333
2     2                0      3.700000
3     3                0      3.833333
4     4                0      2.357143

In [99]: # mean before is computed based on all the users before the beginning of the experiment
mean_before = df2['active_mins'].mean()
var_before = df2['active_mins'].var()
n_before = df2.shape[0]

#compute the mean after rollout but only for users under treatment
mean_after = df[df.variant_number==1]['active_mins'].mean()
var_after = df[df.variant_number==1]['active_mins'].var()
n_after = df[df.variant_number==1].shape[0]

print(n_after)

#print('mean before',mean_before)
#print('mean after',mean_after)

9163

In [101]: # estimated standard deviation of the difference between the two means.
sigma_dif = math.sqrt((var_before/n_before)+(var_after/n_after))
upper = (mean_after-mean_before) + (1.96 * sigma_dif)
lower = (mean_after-mean_before) - (1.96 * sigma_dif)
print([lower,upper])

[1.2698277548728634, 1.372264409083204]

In [103]: #This table contains data about some user attributes.
#Each row represents attributes of a unique user.
t4 = pd.read_csv("t4_user_attributes.csv")
print("the number of records is ",t4.shape)
t4.head()

```

the number of records is (50000, 3)

```
Out[103]:   uid gender  user_type
0      0   male non_reader
1      1   male    reader
2      2   male non_reader
3      3   male non_reader
4      4   male non_reader
```

```
In [104]: t5 = pd.merge(left=t4,right=t2, left_on='uid', right_on='uid')
t5.head()
```

```
Out[104]:   uid gender  user_type  variant_number      dt signup_date
0      0   male non_reader             0 2019-02-06 2018-09-24
1      1   male    reader             0 2019-02-06 2016-11-07
2      2   male non_reader             0 2019-02-06 2018-09-17
3      3   male non_reader             0 2019-02-06 2018-03-04
4      4   male non_reader             0 2019-02-06 2017-03-09
```

```
In [105]: t5.shape
```

```
Out[105]: (50000, 6)
```

```
In [108]: pd.crosstab(t5['variant_number'],t5['gender']).apply(lambda r: r/r.sum(), axis=1)
```

```
Out[108]: gender          female      male  unknown
variant_number
0          0.290175  0.555925   0.1539
1          0.287000  0.552000   0.1610
```

```
In [109]: pd.crosstab(t5['variant_number'],t5['user_type']).apply(lambda r: r/r.sum(), axis=1)
```

```
Out[109]: user_type      contributor  new_user  non_reader    reader
variant_number
0          0.022875  0.091325    0.717475  0.168325
1          0.012900  0.123500    0.736700  0.126900
```

```
In [110]: pd.crosstab(t5['gender'],t5['user_type']).apply(lambda r: r/r.sum(), axis=1)
```

```
Out[110]: user_type  contributor  new_user  non_reader    reader
gender
female          0.017200  0.109898    0.723907  0.148995
male            0.024462  0.083619    0.716108  0.175812
unknown         0.014937  0.125676    0.735127  0.124260
```

```
In [111]: merged_df3 = pd.merge(left=t1,right=t5, left_on='uid', right_on='uid')
merged_df3.head()
```

```
Out[111]:
```

	uid	dt_x	active_mins	gender	user_type	variant_number	\
0	0	2019-02-22	5.0	male	non_reader		0
1	0	2019-03-11	5.0	male	non_reader		0
2	0	2019-03-18	3.0	male	non_reader		0
3	0	2019-03-22	4.0	male	non_reader		0
4	0	2019-04-03	9.0	male	non_reader		0

	dt_y	signup_date
0	2019-02-06	2018-09-24
1	2019-02-06	2018-09-24
2	2019-02-06	2018-09-24
3	2019-02-06	2018-09-24
4	2019-02-06	2018-09-24

```
In [116]: # group by user to compute mean time across all his/her visit.
df3 = merged_df3.groupby(['uid', 'variant_number', 'gender', 'user_type'], as_index=False)
df3.shape # we have 46525 unique users in both groups
#df3.variant_number.value_counts()
```

```
Out[116]: (46525, 5)
```

```
In [117]: df3.head()
```

```
Out[117]:
```

	uid	variant_number	gender	user_type	active_mins
0	0	0	male	non_reader	3.307692
1	1	0	male	reader	9.000000
2	2	0	male	non_reader	2.428571
3	3	0	male	non_reader	3.208333
4	4	0	male	non_reader	1.950000

```
In [171]: df3.to_csv('df3_quora.csv')
```

```
In [122]: #stats = df3.groupby(['variant_number', 'gender', 'user_type'])['active_mins'].agg(['mean', 'count'])
#stats = df3.groupby(['variant_number', 'gender'])['active_mins'].agg(['mean', 'count'])
stats = df3.groupby(['variant_number', 'user_type'])['active_mins'].agg(['mean', 'count'])
```

stats

```
Out[122]:
```

	variant_number	user_type	mean	count	std	var
0	0	contributor	8.067700	885	2.948805	8.695448
1	0	new_user	2.892973	2371	1.828464	3.343280
2	0	non_reader	3.471255	27442	1.679886	2.822016
3	0	reader	6.871204	6664	2.581401	6.663631
4	1	contributor	8.062027	124	2.507654	6.288329
5	1	new_user	4.414679	796	2.522767	6.364354
6	1	non_reader	4.919898	6994	2.034194	4.137944
7	1	reader	7.654943	1249	2.230615	4.975642

```
In [144]: utype='contributor'
control = stats[(stats.variant_number==0) & (stats.user_type==utype)]
```

```
treat = stats[(stats.variant_number==1) & (stats.user_type==utype)]
#sigma_dif = math.sqrt((control['var'][0]/control['count'][0])+((treat['var'][0]/treat
treat.iloc[0]['var']
```

Out[144]: 6.2883294657262665

```
In [153]: #create function to compute confidence interval given a user_type category and data fr
def ci_usertype(stats,utype):
    control = stats[(stats.variant_number==0) & (stats.user_type==utype)]
    treat = stats[(stats.variant_number==1) & (stats.user_type==utype)]

    sigma_dif = math.sqrt((control.iloc[0]['var']/control.iloc[0]['count'])+((treat.il
    upper = (treat.iloc[0]['mean']-control.iloc[0]['mean']) + (1.96 * sigma_dif)
    lower = (treat.iloc[0]['mean']-control.iloc[0]['mean']) - (1.96 * sigma_dif)
    return (lower,upper)
```

(-0.4879189443104225, 0.47657393324998065)

```
In [155]: # run the function on all unique user_type categories
```

```
utypes=stats['user_type'].unique()
utypes=utypes.tolist()
utypes

for u in utypes:
    inter = ci_usertype(stats,u)
    print(u, '--->', inter)
```

```
contributor ---> (-0.4879189443104225, 0.47657393324998065)
new_user ---> (1.3316220606045999, 1.7117911439585125)
non_reader ---> (1.3969912190555374, 1.5002947542414995)
reader ---> (0.6453729886947472, 0.9221051371502116)
```

```
In [156]: stats2 = df3.groupby(['variant_number','gender'])['active_mins'].agg(['mean', 'count',
stats2
```

```
Out[156]:
```

	variant_number	gender	mean	count	std	var
0	0	female	3.961722	10725	2.305638	5.315969
1	0	male	4.318368	20961	2.487382	6.187067
2	0	unknown	3.883059	5676	2.248822	5.057200
3	1	female	5.147343	2599	2.392395	5.723552
4	1	male	5.403433	5112	2.318357	5.374780
5	1	unknown	5.154455	1452	2.348080	5.513481

```

In [158]: #create function to compute confidence interval given a user_type category and data fr
def ci_gender(stats,gen):
    control = stats[(stats.variant_number==0) & (stats.gender==gen)]
    treat = stats[(stats.variant_number==1) & (stats.gender==gen)]

    sigma_dif = math.sqrt((control.iloc[0]['var']/control.iloc[0]['count'])+((treat.il
    upper = (treat.iloc[0]['mean']-control.iloc[0]['mean']) + (1.96 * sigma_dif)
    lower = (treat.iloc[0]['mean']-control.iloc[0]['mean']) - (1.96 * sigma_dif)
    return (lower,upper)

In [161]: ugender=stats2['gender'].unique()
ugender=ugender.tolist()
ugender

for g in ugender:
    inter = ci_gender(stats2,g)
    print(g, '--->', inter)

female ---> (1.0838168999869435, 1.2874258861290009)
male ---> (1.0131414147714226, 1.1569885842349406)
unknown ---> (1.1371945853770327, 1.405597029108042)

```

```

In [163]: stats3 = df3.groupby(['variant_number','gender','user_type'])['active_mins'].agg(['mea
stats3

```

```

Out[163]:
   variant_number  gender  user_type    mean  count    std    var
0              0  female  contributor  7.142652   220  2.676451  7.163392
1              0  female    new_user  2.822876   762  1.777076  3.157999
2              0  female  non_reader  3.374448  7935  1.700659  2.892242
3              0  female    reader  6.632086  1808  2.494685  6.223454
4              0   male  contributor  8.492587   572  2.924966  8.555426
5              0   male    new_user  2.955687  1132  1.840741  3.388328
6              0   male  non_reader  3.537730 15178  1.666697  2.777879
7              0   male    reader  7.015949  4079  2.611095  6.817815
8              0  unknown  contributor  7.642702    93  3.170050 10.049217
9              0  unknown    new_user  2.856120   477  1.878277  3.527926
10             0  unknown  non_reader  3.415631  4329  1.677374  2.813585
11             0  unknown    reader  6.667744   777  2.571137  6.610747
12             1  female  contributor  7.844039    26  2.208249  4.876364
13             1  female    new_user  4.464382   274  2.587578  6.695559
14             1  female  non_reader  4.810470  1968  2.142444  4.590068
15             1  female    reader  7.503789   331  2.141313  4.585221
16             1   male  contributor  8.320709    79  2.506627  6.283178
17             1   male    new_user  4.354160   369  2.472309  6.112312
18             1   male  non_reader  4.998772  3920  1.973386  3.894253
19             1   male    reader  7.746163   744  2.221167  4.933583
20             1  unknown  contributor  7.284756    19  2.823186  7.970380
21             1  unknown    new_user  4.471628   153  2.539061  6.446829

```


22	1	unknown	non_reader	4.835059	1106	2.038206	4.154286
23	1	unknown	reader	7.552438	174	2.422491	5.868462

```
In [167]: covariates = [(a,b) for a in ugender for b in utypes]
covariates
len(covariates)
```

```
Out[167]: 12
```

```
In [169]: def ci_covariates(stats,cov):
    control = stats[(stats.variant_number==0) & (stats.gender==cov[0]) & (stats.user_type==cov[1])]
    treat = stats[(stats.variant_number==1) & (stats.gender==cov[0]) & (stats.user_type==cov[1])]

    sigma_dif = math.sqrt((control.iloc[0]['var']/control.iloc[0]['count'])+((treat.iloc[0]['var']/treat.iloc[0]['count'])))
    upper = (treat.iloc[0]['mean']-control.iloc[0]['mean']) + (1.96 * sigma_dif)
    lower = (treat.iloc[0]['mean']-control.iloc[0]['mean']) - (1.96 * sigma_dif)
    return (lower,upper)
```

```
In [170]: for c in covariates:
    inter = ci_covariates(stats3,c)
    print(c, '--->', inter)
```

```
('female', 'contributor') ---> (-0.2181710324317161, 1.6209454206671288)
('female', 'new_user') ---> (1.3101518625750066, 1.9728603216581513)
('female', 'non_reader') ---> (1.3342376926849375, 1.5378077703842925)
('female', 'reader') ---> (0.613943688075681, 1.1294617876307897)
('male', 'contributor') ---> (-0.7743697022060778, 0.43061390495753704)
('male', 'new_user') ---> (1.1243689244004251, 1.6725771584888394)
('male', 'non_reader') ---> (1.3938146345782447, 1.5282684611886996)
('male', 'reader') ---> (0.5516221399068965, 0.9088071350362895)
('unknown', 'contributor') ---> (-1.781544948687302, 1.0656539264384974)
('unknown', 'new_user') ---> (1.179293254438179, 2.0517219238093998)
('unknown', 'non_reader') ---> (1.2893264050791, 1.5495292036690422)
('unknown', 'reader') ---> (0.48189251181388243, 1.2874962369534175)
```

```
In [ ]:
```

```

library(readr)
library(ggplot2)
df3 <- read_csv("quora_challenge/df3_quora.csv",
                 col_types = cols(X1 = col_skip()))
View(df3)

nrow(df3)

table(df3$variant_number,exclude=NULL)

#boxplot user type
g <- ggplot(df3, aes(user_type, active_mins))
g + geom_boxplot(aes(fill=factor(variant_number)))+
  scale_fill_manual(values=c('#107896', '#AA2200'))+
  theme(
    panel.background = element_blank(),
    panel.grid.major = element_line(colour = "grey",size=0.01,linetype =
8),
    panel.grid.minor = element_blank(),
    axis.line.x = element_line(colour = "black"),
    axis.line.y = element_line(colour = "black"),
    #legend.background = element_rect(fill="transparent", size=0.5,
linetype="solid"),
    #legend.position = c(0.1,0.9),
    legend.text = element_text(colour="black", size=14),
    legend.title=element_text(size=14),
    legend.box = "horizontal",
    axis.text = element_text(colour = "black"),
    axis.text.x = element_text(size = 14,vjust=0.6),
    axis.text.y = element_text(size = 14),
    axis.title = element_text(size = 14),
    strip.text.x = element_text(size = 14)) + labs(x='User Type',y='Active
Mins') + guides(fill=guide_legend(title="Treatment"))

#boxplot user gender
g <- ggplot(df3, aes(gender, active_mins))
g + geom_boxplot(aes(fill=factor(variant_number)))+
  scale_fill_manual(values=c('#107896', '#AA2200'))+
  theme(
    panel.background = element_blank(),
    panel.grid.major = element_line(colour = "grey",size=0.01,linetype =
8),
    panel.grid.minor = element_blank(),
    axis.line.x = element_line(colour = "black"),
    axis.line.y = element_line(colour = "black"),
    #legend.background = element_rect(fill="transparent", size=0.5,
linetype="solid"),
    #legend.position = c(0.1,0.9),
    legend.text = element_text(colour="black", size=14),
    legend.title=element_text(size=14),
    legend.box = "horizontal",
    axis.text = element_text(colour = "black"),

```

```

axis.text.x = element_text(size = 14,vjust=0.6),
axis.text.y = element_text(size = 14),
axis.title = element_text(size = 14),
strip.text.x = element_text(size = 14)) + labs(x='User
Gender',y='Active Mins') + guides(fill=guide_legend(title="Treatment"))

#boxplot interaction

df3$f1f2 <- interaction(df3$user_type, df3$gender)

g <- ggplot(df3, aes(f1f2, active_mins))
g + geom_boxplot(aes(fill=factor(variant_number)))+
  scale_fill_manual(values=c('#107896','#AA2200'))+
  theme(
    panel.background = element_blank(),
    panel.grid.major = element_line(colour = "grey",size=0.01,linetype =
8),
    panel.grid.minor = element_blank(),
    axis.line.x = element_line(colour = "black"),
    axis.line.y = element_line(colour = "black"),
    #legend.background = element_rect(fill="transparent", size=0.5,
linetype="solid"),
    #legend.position = c(0.1,0.9),
    legend.text = element_text(colour="black", size=14),
    legend.title=element_text(size=14),
    legend.box = "horizontal",
    axis.text = element_text(colour = "black"),
    axis.text.x = element_text(size = 14,vjust=0.6,angle=90),
    axis.text.y = element_text(size = 14),
    axis.title = element_text(size = 14),
    strip.text.x = element_text(size = 14)) + labs(x='User Type-
Gender',y='Active Mins') + guides(fill=guide_legend(title="Treatment"))

```